

# Hierarchical Agglomerative Clustering

---

CHIEN CHIN CHEN

# Preface

---

Different to flat clustering, **hierarchical clustering** outputs a **hierarchy**.

- A structure (tree) that is supposed to be more informative than the unstructured set of clusters in flat clustering.
- Some researchers believe that hierarchical clustering produces better clusters than flat clustering But ... there is no consensus on this issue.

This chapter focuses on **hierarchical agglomerative clustering** (HAC).

- Present five different linkage measures: single-link, complete-link, group-average, centroid similarity, and Ward.
- Clustering result is **deterministic**!!

# Hierarchical Agglomerative Clustering (1/6)

---

A **bottom-up** approach.

Treats each document as a singleton cluster at the outset.

**Successively merge** (or agglomerate) pairs of clusters until all clusters have been merged into a single cluster.

An HAC clustering is typically visualized as a ***dendrogram***.

# Hierarchical Agglomerative Clustering (2/6)


A merge of two clusters is represented as a horizontal line.

The y-axis represents combination similarity (or distance).

- The similarity of the two clusters connected by the horizontal line.

The y-axis at leaf nodes is 1.

- To represent each cluster as a document.



# Hierarchical Agglomerative Clustering (3/6)

SimpleHAC ( $d_1, \dots, d_N$ )

```
for  $n \leftarrow 1$  to  $N$ 
do for  $i \leftarrow 1$  to  $N$ 
  do  $C[n][i] \leftarrow \text{Sim}(d_n, d_i)$ 
   $I[n] \leftarrow 1$ 
 $A \leftarrow []$ 
```

```
for  $k \leftarrow 1$  to  $N-1$ 
  do  $\langle i, m \rangle \leftarrow \text{argmax}_{\{ \langle i, m \rangle : i \neq m \text{ and } I[i]=1 \text{ and } I[m]=1 \}} C[i][m]$ 
   $A.\text{Append}(\langle i, m \rangle)$ 
  for  $j \leftarrow 1$  to  $N$ 
    do  $C[i][j] \leftarrow \text{Sim}(j, i, m)$ 
        $C[j][i] \leftarrow \text{Sim}(j, i, m)$ 
        $I[m] \leftarrow 0$ 
```

return  $A$


$C[i][j]$ : the similarity between clusters  $i$  and  $j$ .

$I$ : indicate which clusters are still available to be merged.

$A$ : a list of merges

the similarity of cluster  $j$  with the merge of cluster  $i$  and  $m$ .


# Hierarchical Agglomerative Clustering (4/6)



$$I = \begin{array}{c|cccccccccc} & I[1] & I[2] & I[3] & I[4] & I[5] & I[6] & I[7] & I[8] & I[9] & I[10] \\ \hline & 1 & 0 & 1 & 1 & 1 & 1 & 1 & 1 & 0 & 1 \end{array}$$

$$A = \{ \langle 3,9 \rangle \quad \langle 1,2 \rangle \quad \dots \}$$

1. Search for the max similarity pair -  $\langle i, m \rangle$
2. Record and merge cluster  $m$  into cluster  $i$ .
3. Update the pair similarities between the new merged cluster  $i$  and other clusters.
4.  $I[m] = 0$ , remove the merged cluster  $m$ .



So complicated ...  
No worry, SKLearn helps us do it in an easy way!!

# Hierarchical Agglomerative Clustering (5/6)

---

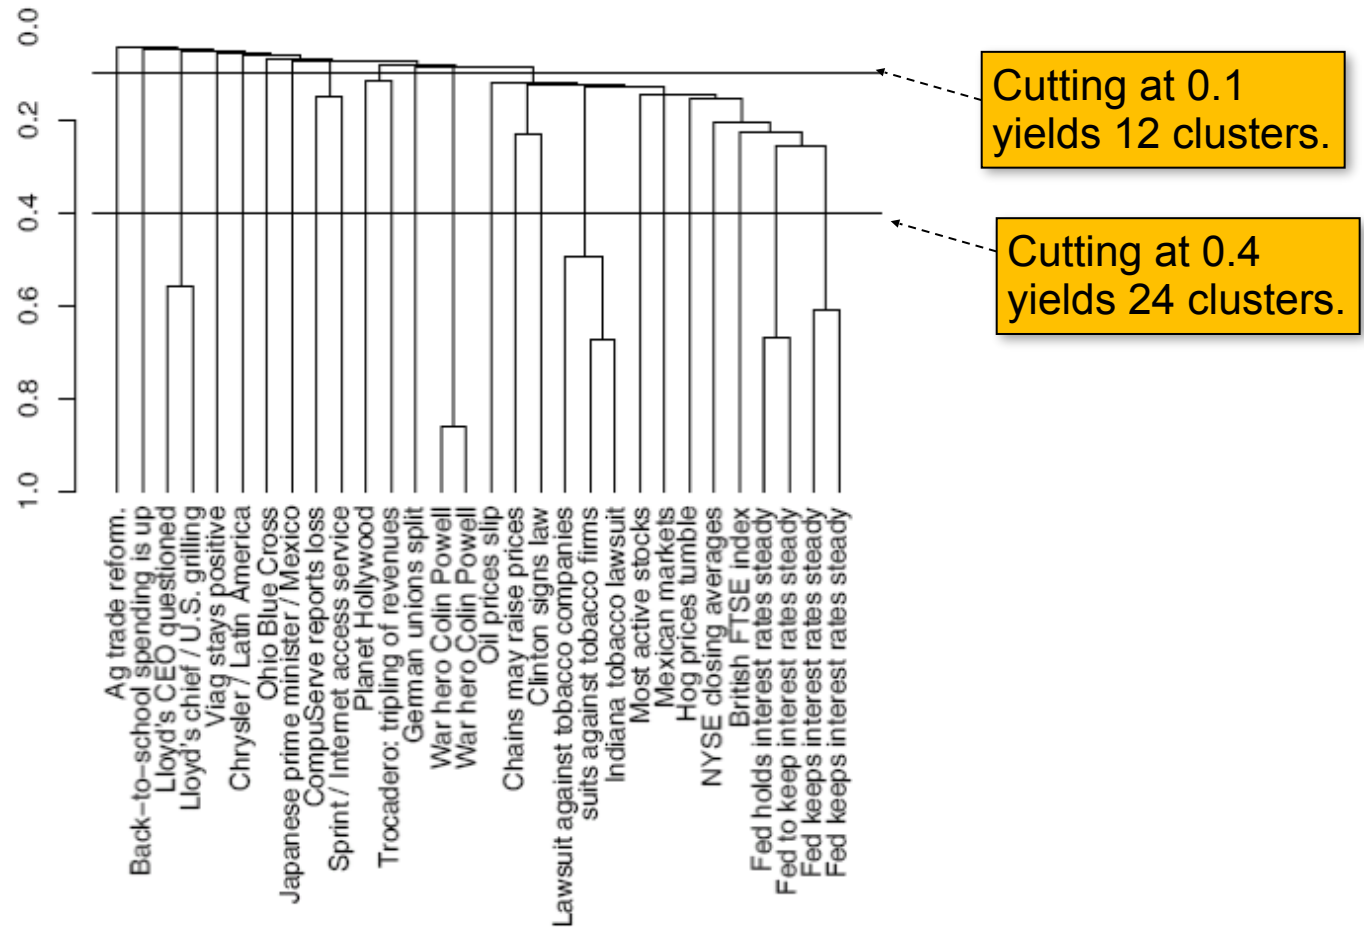
In some cases, we want a partition of disjoint clusters just as in flat clustering.

- The hierarchy needs to be cut at some point!!

## **Criteria of cutting:**

- As in flat clustering, we can also pre-specify the number of clusters  $K$ .
- Cut at a pre-specified level of similarity.
  - We cut the dendrogram at 0.4.
  - We want clusters with a minimum combination similarity of 0.4.
- Cut the dendrogram where the gap between two successive combination similarities is largest.
  - Merging one more cluster decreases the quality of the clustering significantly.

# Hierarchical Agglomerative Clustering (6/6)






# Single-link and Complete-link (1/7)

## *Single-link :*

- The similarity of two clusters is the similarity of **their most similar members**.
- This merge criterion is **local**.
  - We pay attention solely to the area where the two clusters come closest to each other.
  - The cluster overall structure is not taken into account.




(a) single link: maximum similarity

# Single-link and Complete-link (2/7)


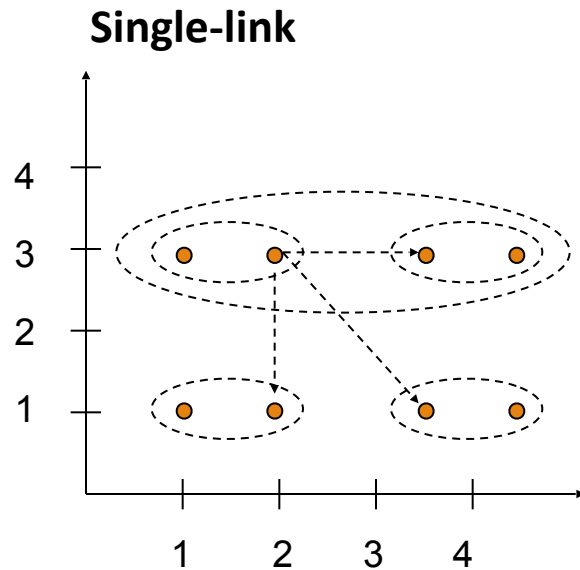
## **Complete-link :**

- The similarity of two cluster is the similarity of **their most dissimilar members**.
- This merge criterion is **non-local**.
  - The entire structure of the clustering can influence merge decision.
- Is equivalent to choosing the cluster pair whose merge has the smallest diameter.



# Single-link and Complete-link (3/7)

---



# Single-link and Complete-link (4/7)

---

Graph-theoretic interpretations:

- Let  $s_k$  to be the combination similarity of the two clusters merged in step  $k$ .
- $G(s_k)$  the graph that links all data points with a similarity of at least  $s_k$ .
- Then, the clusters after step  $k$  in **single-link** clustering are the **connected components** of  $G(s_k)$ .
  - Points in a clusters are linked via at least one link → **single-link!!**
- And the clusters after step  $k$  in **complete-link** clustering are maximal **cliques** of  $G(s_k)$ .
  - A clique is a set of points that are completely linked with each other → **complete-link!!**


# Single-link and Complete-link (5/7)

Single-link assesses cluster quality via a pair of document.

- The two most similar documents.
- Clustering based on one pair cannot fully reflect the distribution of documents in a cluster.


**Chaining:** a problem of single-link.

- A chain of points can be extended for long distances without regard to the overall shape of the cluster.




► **Figure 17.6** Chaining in single-link clustering. The local criterion in single-link clustering can cause undesirable elongated clusters.

# Single-link and Complete-link (6/7)



► **Figure 17.1** A dendrogram of a single-link clustering of 30 documents from Reuters-RCV1. The y-axis represents combination similarity, the similarity of the



► **Figure 17.5** A dendrogram of a complete-link clustering of the 30 documents in Figure 17.1.

The last 12 merges add on  
single documents → chaining!!


We **often** prefer compact clusters with  
small diameters over long, straggly clusters.

# Single-link and Complete-link (7/7)

---

The problem of complete-link clustering – **very sensitive to outliers**.

- A single document far from the center can increase diameters of candidate merge clusters dramatically.



# Group-average Linkage (1/2)

## **Group-average:**

- The similarity of two cluster is the average similarity of all pairs of documents.
- **Including pairs from the same cluster!!**
- **But self-similarities are not included in the average!!**

$$sim - ga(\omega_i, \omega_j) = \frac{1}{(N_i + N_j)(N_i + N_j - 1)} \sum_{d_m \in \omega_i} \sum_{d_n \in \omega_j, d_n \neq d_m} d_m \cdot d_n$$

clusters  $i$  and  $j$

the number of documents in cluster  $i$

the **normalized** document vectors



# Group-average Linkage (2/2)

---

Why not including self-similarity?

- If we define group-average similarity as including self-similarities:

$$sim - ga'(\omega_i, \omega_j) = \frac{1}{(N_i + N_j)^2} \sum_{d_m \in \omega_i \cup \omega_j} \sum_{d_n \in \omega_i \cup \omega_j} \vec{d}_m \cdot \vec{d}_n$$

- For a cluster of size  $i$ , the proportion of self-similarities is  $i/i^2$  or  $1/i$ .
  - **This gives an unfair advantage to small clusters.**
- Moreover, for two documents  $d_1, d_2$  with similarity  $s$ 
  - $sim-ga'(d_1, d_2) = (2+2s)/4 = (1+s)/2$ .
  - $sim-ga(d_1, d_2) = 2s/2 = s$ .
  - is the same as in single-link, complete-link, and centroid clustering.

# Centroid Linkage

---

## **Centroid linkage:**

- The similarity of two cluster is defined as the similarity of their centroids:

$$\begin{aligned} \text{sim} - \text{cent}(\omega_i, \omega_j) &= \left( \frac{1}{N_i} \sum_{d_m \in \omega_i} \vec{d}_m \right) \cdot \left( \frac{1}{N_j} \sum_{d_n \in \omega_j} \vec{d}_n \right) \\ &= \frac{1}{N_i N_j} \sum_{d_m \in \omega_i} \sum_{d_n \in \omega_j} \vec{d}_m \cdot \vec{d}_n \end{aligned}$$

- The similarity is equivalent to average similarity of all pairs of documents from **different** clusters.

# Ward's Linkage (1/2)

---

## **Ward's Linkage:**

- The distance between two clusters is how much the sum of squares will increase when we merge them:

$$Ward(\omega_i, \omega_j) = \sum_{d_m \in \omega_i \cup \omega_j} \left\| \vec{d}_m - centroid_{\omega_i \cup \omega_j} \right\|^2 -$$

The sum of squares starts out at zero, and grows as we merge clusters


$$\sum_{d_m \in \omega_i} \left\| \vec{d}_m - centroid_{\omega_i} \right\|^2 + \sum_{d_m \in \omega_j} \left\| \vec{d}_m - centroid_{\omega_j} \right\|^2$$

- In the beginning, every document (vector) is in its own cluster.


Ward's linkage keeps this growth as small as possible.

# Ward's Linkage (2/2)

---



Which two clusters would Ward's linkage group first?



Another interpretation of Ward's linkage is to minimize the variance of the clusters being merged.

# Which Similarity Strategy Is The Best ?

---

[Voorhees 1985] recommends complete-link and centroid clustering over single-link for a retrieval application.

But in news-related applications (e.g., event/topic detection), single-link seems to be a better strategy.

- Allan et al. (1998) apply the single-link clustering to first story detection.

# Practice Time (1/2)

---

```
In [8]: from sklearn.cluster import AgglomerativeClustering
        hac = AgglomerativeClustering(linkage='ward', n_clusters=2)
        hac.fit(TFIDF_vectors.toarray())
```

```
Out[8]: AgglomerativeClustering(affinity='euclidean', compute_full_tree='auto',
                                connectivity=None, distance_threshold=None,
                                linkage='ward', memory=None, n_clusters=2)
```

```
In [9]: hac.labels_
```

```
Out[9]: array([0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0])
```

## Parameters:

- linkage: ward, single, complete, average
  - default – ward
- n\_clusters: the number of clusters to find.
- distance\_threshold: the distance threshold to stop merging

# Practice Time (2/2)

## How to show the cluster hierarchy?

```
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from scipy.cluster.hierarchy import dendrogram
```

```
def plot_dendrogram(model, **kwargs):
    # https://scikit-learn.org/stable/auto\_examples/cluster/plot\_agglomerative\_dendrogram.html#sphx-glr-auto
    # Children of hierarchical clustering
    children = model.children_

    # Distances between each pair of children
    # Since we don't have this information, we can use a uniform one for plotting
    distance = np.arange(children.shape[0])

    # The number of observations contained in each cluster level
    no_of_observations = np.arange(2, children.shape[0]+2)

    # Create linkage matrix and then plot the dendrogram
    linkage_matrix = np.column_stack([children, distance, no_of_observations]).astype(float)

    # Plot the corresponding dendrogram
    dendrogram(linkage_matrix, **kwargs)
```

```
In [12]: plt.title('Hierarchical Clustering Dendrogram')
plot_dendrogram(hac, labels=hac.labels_)
plt.show()
```

